

Comparative Analysis — Our v0 Options

Version: 2025-12-25 • Purpose: pick the sharpest v0 wedge and sequencing.

The candidates (what each v0 is)

Below are the main v0 directions we've discussed across the project. They are not mutually exclusive; the goal is to pick **the smallest product that proves the core wedge**, then sequence the rest.

V0-A — Agent Mesh Observatory

Observability-first: capture agent/tool graph, costs, latency, loops; dashboards + kill switches.
Primary buyer/user: AI platform / infra teams; devex leaders; SRE-style buyers.

V0-B — Shadow AI Runtime Broker

Workflow interception: browser/IDE + local proxy + LLM gateway reroute to sanctioned endpoints; redaction and policy at point of use.
Primary buyer/user: Security + IT (SSE/DLP overlap) and DevEx for IDE coverage.

V0-C — Agent Identity Control Plane

Action-bound governance: intent-bound capabilities for tool calls + credential brokering + enforce at tool/MCP boundary; generate action receipts.
Primary buyer/user: Security/Platform teams enabling agents safely; SOC/IR visibility.

V0-D — Agent Flight Recorder (IR/Forensics)

Post-incident evidence: timeline, graph snapshots, replayable traces, exportable evidence bundles; less enforcement, more investigation.
Primary buyer/user: Security operations + platform reliability; governance programs.

V0-E — PR-Centric Voice Manager (Dogfood Workflow)

A concrete autonomous workflow used to dogfood: PR creation/review orchestration + voice/command interface; proves integration and receipts.
Primary buyer/user: Internal productivity + demos; not the main product wedge.

Comparison matrix (qualitative)

v0 option	Core primitive	Fastest proof of value	Main risk / failure mode	Moat lever	What we'd build first
V0-A Observatory	Unified agent/tool graph + cost/latency signals	Show a live topology + catch a loop/cost storm + kill switch works	Looks like “observability tooling” (crowded) unless tied to enforcement/IR	Graph schema + cross-runtime instrumentation + anomaly models	SDK/proxy ingest + graph UI + loop detector + node-level gating
V0-B Shadow AI Broker	Workflow capture + reroute to sanctioned AI + pre-exfil checks	Reduce unsanctioned GenAI usage while preserving speed	Crowded vs SSE/browser-DLP; hard to cover all clients; policy UX	Endpoint coverage + “approved path” UX + enterprise distribution	Browser ext + IDE plugin + gateway base_url swap + redaction

V0-C Identity Control Plane	Intent-bound capabilities + tool-action enforcement + receipts	Safely enable agents to act (write/modify) without handing raw creds	Integration complexity; needs 2–3 killer tool adapters to feel real	Capability model + receipts + cross-runtime normalization	Capability issuer + tool proxy (GitHub/Jira/Slack) + receipt ledger
V0-D Flight Recorder	Receipt ledger + trace graph + evidence bundles	IR teams can answer “what happened” quickly during incidents	“Nice to have” without incidents; must show cost/time saved in investigations	Evidence standard + replay/forensics depth + SIEM workflows	Append-only ledger + timeline UI + export bundles + replay stub
V0-E PR Voice Manager	Autonomous workflow orchestration + repo safety rails	Compelling demo + dogfood data; showcases receipts/controls	Becomes a side-product; can distract from core buyer/budget	Demonstrates end-to-end; generates ground-truth action traces	One repo workflow + constrained tool set + approvals + logging

Decision scoring (1–5) — what to ship first

Scores are directional (not “truth”). They’re here to force tradeoffs and sequencing.

v0	Pain	Pay	Crowding	Time	Diff	Platform	Total
V0-A Observatory	3	3	2	3	3	4	18
V0-B Shadow AI Broker	4	4	2	3	3	4	20
V0-C Identity Control Plane	5	5	4	3	5	5	27
V0-D Flight Recorder	4	4	4	4	4	5	25
V0-E PR Voice Manager	2	2	5	4	2	2	17

Recommendation: C → D → A (with E as dogfood)

Ship V0-C (Identity Control Plane) as the core wedge: capability issuance + tool-action enforcement + receipts. Then **add V0-D (Flight Recorder)** as the first “SOC value” expansion (evidence bundles and replay). Once enforcement + receipts exist, the **Observatory (V0-A)** becomes a natural UI layer instead of a standalone observability tool. Use **V0-E** (PR workflow) to dogfood and generate compelling demos and traces. Treat **V0-B** (shadow AI broker) as a future distribution channel if needed, not the initial battlefield.

What “V0-C done” means (minimal but real)

- One capability issuer (JWT-like) with ABAC claims + TTL + reason codes.
- Two tool adapters (pick: GitHub + Jira) enforcing capability checks per action.
- Credential broker (short-lived grants; never hand raw tokens to the model).
- Receipt ledger (append-only, hash-chained) + basic query UI (who did what, when, why).
- Demo: agent attempts disallowed action → blocked with rationale; allowed action → executed with receipt.

Sequencing notes (why not start elsewhere)

Why not V0-A first? Without enforcement/receipts, observability alone risks being perceived as “another dashboard.”

Why not V0-B first? Shadow AI controls are crowded and distribution-heavy; it's a strong feature/add-on after we own tool-action governance.

Why V0-D second? It monetizes for security programs and makes every deployment stickier via evidence and workflows.

Why keep V0-E? It's the fastest way to produce realistic traces and demos without waiting for design partners.

One-page technical sketch: Control Plane + Enforcement + Evidence

CONTROL PLANE

- Policy service (ABAC) + approvals
- Capability issuer (mint scoped tokens)
- Registry: agents, tools, resources, environments

ENFORCEMENT

- Tool proxy / MCP gateway adapter
- OAuth/secret broker (short-lived grants)
- Parameter constraints + data minimization

EVIDENCE

- Receipt ledger (hash-chained)
- Trace graph: request -> tool calls -> side effects
- Export: incident evidence bundle (JSON + human PDF)